


# Principles of High Performance Computing (ICS 632)

## Algorithms on a Grid of Processors (II)

### 2-D Matrix Distribution



- We denote by  $a_{ij}$  an element of the matrix
- We denote by  $A_{ij}$  (or  $A_{ij}$ ) the block of the matrix allocated to  $P_{i,j}$

$C_{00}$	$C_{01}$	$C_{02}$	$C_{03}$
$C_{10}$	$C_{11}$	$C_{12}$	$C_{13}$
$C_{20}$	$C_{21}$	$C_{22}$	$C_{23}$
$C_{30}$	$C_{31}$	$C_{32}$	$C_{33}$

$A_{00}$	$A_{01}$	$A_{02}$	$A_{03}$
$A_{10}$	$A_{11}$	$A_{12}$	$A_{13}$
$A_{20}$	$A_{21}$	$A_{22}$	$A_{23}$
$A_{30}$	$A_{31}$	$A_{32}$	$A_{33}$

$B_{00}$	$B_{01}$	$B_{02}$	$B_{03}$
$B_{10}$	$B_{11}$	$B_{12}$	$B_{13}$
$B_{20}$	$B_{21}$	$B_{22}$	$B_{23}$
$B_{30}$	$B_{31}$	$B_{32}$	$B_{33}$

### The Cannon Algorithm

- This is a very old algorithm
  - From the time of systolic arrays
  - Adapted to a 2-D grid
- The algorithm starts with a redistribution of matrices A and B
  - Called “preskewing”
- Then the matrices are multiplied
- Then the algorithms are re-redistributed to match the initial distribution
  - Called “postskewing”

### Cannon’s Preskewing

- Matrix A: each block row of matrix A is shifted so that each processor in the first processor column holds a diagonal block of the matrix

$A_{00}$	$A_{01}$	$A_{02}$	$A_{03}$
$A_{10}$	$A_{11}$	$A_{12}$	$A_{13}$
$A_{20}$	$A_{21}$	$A_{22}$	$A_{23}$
$A_{30}$	$A_{31}$	$A_{32}$	$A_{33}$

$A_{00}$	$A_{01}$	$A_{02}$	$A_{03}$
$A_{11}$	$A_{12}$	$A_{13}$	$A_{14}$
$A_{22}$	$A_{23}$	$A_{20}$	$A_{21}$
$A_{33}$	$A_{30}$	$A_{31}$	$A_{32}$

### Cannon’s Preskewing

- Matrix B: each block column of matrix B is shifted so that each processor in the first processor row holds a diagonal block of the matrix

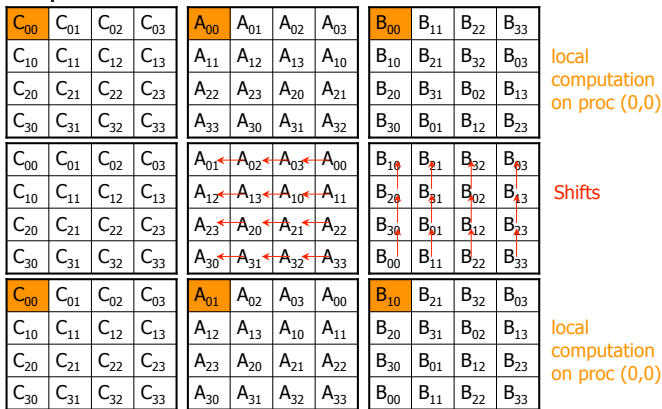
$B_{00}$	$B_{01}$	$B_{02}$	$B_{03}$
$B_{10}$	$B_{11}$	$B_{12}$	$B_{13}$
$B_{20}$	$B_{21}$	$B_{22}$	$B_{23}$
$B_{30}$	$B_{31}$	$B_{32}$	$B_{33}$

$B_{00}$	$B_{11}$	$B_{22}$	$B_{33}$
$B_{10}$	$B_{21}$	$B_{32}$	$B_{03}$
$B_{20}$	$B_{31}$	$B_{02}$	$B_{13}$
$B_{30}$	$B_{01}$	$B_{12}$	$B_{23}$

### Cannon’s Computation

- The algorithm proceeds in  $q$  steps
- At each step each processor performs the multiplication of its block of A and adds the result to its block of C
- Then blocks of A are shifted to the left and blocks of B are shifted upward
  - Blocks of C never move
- Let’s see it on a picture

## Cannon's Steps



## The Algorithm

Participate in preskewing of A  
 Participate in preskewing of B  
 For  $k = 1$  to  $q$   
     Local  $C = C + A * B$   
     Vertical shift of B  
     Horizontal shift of A  
 Participate in postskewing of A  
 Participate in postskewing of B

## Performance Analysis

- Let's do a simple performance analysis with a 4-port model
  - The 1-port model is typically more complicated
- Symbols
  - $n$ : size of the matrix
  - $q \times q$ : size of the processor grid
  - $m = n / q$
  - $b$ : communication start-up cost
  - $T_a$ : time to do a basic computation ( $+$ ,  $-$ ,  $*$ ,  $/$ )
  - $T_c$ : time to communicate a matrix element
- $T(m, q) = T_{\text{preskew}} + T_{\text{compute}} + T_{\text{postskew}}$

## Pre/Post-skewing times

- Let's consider the horizontal shift
- Each row must be shifted so that the diagonal block ends up on the first column
- On a mono-directional ring:
  - The last row needs to be shifted  $(q-1)$  times
  - All rows can be shifted in parallel
  - Total time needed:  $(q-1)(b + m^2 T_c)$
- On a bi-directional ring, a row can be shifted left or right, depending on which way is shortest!
  - A row is shifted at most  $\text{floor}(q/2)$  times
  - All rows can be shifted in parallel
  - Total time needed:  $\text{floor}(q/2)(b + m^2 T_c)$
- Because of the 4-port assumption, preskewing of A and B can occur in parallel (horizontal and vertical shifts do not interfere)
- Therefore:  $T_{\text{preskew}} = T_{\text{postskew}} = \text{floor}(q/2)(b + m^2 T_c)$

## Time for each step

- At each step, each processor computes an  $m \times m$  matrix multiplication
  - Compute time:  $m^3 T_a$
- At each step, each processor sends/receives a  $m \times m$  block in its processor row and its processor column
  - Both can occur simultaneously with a 4-port model
  - Takes time  $b + m^2 T_c$
- Therefore, the total time for the  $q$  steps is:  
 $T_{\text{compute}} = q \max(b + m^2 T_c, m^3 T_a)$

## Cannon Performance Model

- $T(m, n) = 2 * \text{floor}(q/2)(b + m^2 T_c) + q \max(m^3 T_a, b + m^2 T_c)$
- This performance model is easily adapted
  - If one assumes mono-directional links, then the " $\text{floor}(q/2)$ " above becomes " $(q-1)$ "
  - If one assumes 1-port, there is a factor 2 added in front of communication terms
  - If one assumes no overlap of communication and computation at a processor, the " $\max$ " above becomes a sum

## The Fox Algorithm

- This algorithm was originally developed to run on a hypercube topology
  - But in fact it uses a grid, embedded in the hypercube
- This algorithm requires no pre- or post-skewing
- It relies on horizontal broadcasts of the diagonals of matrix A and on vertical shifts of matrix B
- Sometimes called the “multiply-broadcast-roll” algorithm
- Let’s see it on a picture
  - Although it’s a bit awkward to draw because of the broadcasts

## Execution Steps...

C <sub>00</sub>	C <sub>01</sub>	C <sub>02</sub>	C <sub>03</sub>	A <sub>00</sub>	A <sub>01</sub>	A <sub>02</sub>	A <sub>03</sub>	B <sub>00</sub>	B <sub>01</sub>	B <sub>02</sub>	B <sub>03</sub>	initial state
C <sub>10</sub>	C <sub>11</sub>	C <sub>12</sub>	C <sub>13</sub>	A <sub>10</sub>	A <sub>11</sub>	A <sub>12</sub>	A <sub>13</sub>	B <sub>10</sub>	B <sub>11</sub>	B <sub>12</sub>	B <sub>13</sub>	
C <sub>20</sub>	C <sub>21</sub>	C <sub>22</sub>	C <sub>23</sub>	A <sub>20</sub>	A <sub>21</sub>	A <sub>22</sub>	A <sub>23</sub>	B <sub>20</sub>	B <sub>21</sub>	B <sub>22</sub>	B <sub>23</sub>	
C <sub>30</sub>	C <sub>31</sub>	C <sub>32</sub>	C <sub>33</sub>	A <sub>30</sub>	A <sub>31</sub>	A <sub>32</sub>	A <sub>33</sub>	B <sub>30</sub>	B <sub>31</sub>	B <sub>32</sub>	B <sub>33</sub>	
C <sub>00</sub>	C <sub>01</sub>	C <sub>02</sub>	C <sub>03</sub>	A <sub>00</sub>	A <sub>00</sub>	A <sub>00</sub>	A <sub>00</sub>	B <sub>00</sub>	B <sub>01</sub>	B <sub>02</sub>	B <sub>03</sub>	Broadcast of A's 1st diag. (stored in a Separate buffer)
C <sub>10</sub>	C <sub>11</sub>	C <sub>12</sub>	C <sub>13</sub>	A <sub>11</sub>	A <sub>11</sub>	A <sub>11</sub>	A <sub>11</sub>	B <sub>10</sub>	B <sub>11</sub>	B <sub>12</sub>	B <sub>13</sub>	
C <sub>20</sub>	C <sub>21</sub>	C <sub>22</sub>	C <sub>23</sub>	A <sub>22</sub>	A <sub>22</sub>	A <sub>22</sub>	A <sub>22</sub>	B <sub>20</sub>	B <sub>21</sub>	B <sub>22</sub>	B <sub>23</sub>	
C <sub>30</sub>	C <sub>31</sub>	C <sub>32</sub>	C <sub>33</sub>	A <sub>33</sub>	A <sub>33</sub>	A <sub>33</sub>	A <sub>33</sub>	B <sub>30</sub>	B <sub>31</sub>	B <sub>32</sub>	B <sub>33</sub>	
C <sub>00</sub>	C <sub>01</sub>	C <sub>02</sub>	C <sub>03</sub>	A <sub>00</sub>	A <sub>00</sub>	A <sub>00</sub>	A <sub>00</sub>	B <sub>00</sub>	B <sub>01</sub>	B <sub>02</sub>	B <sub>03</sub>	Local computation
C <sub>10</sub>	C <sub>11</sub>	C <sub>12</sub>	C <sub>13</sub>	A <sub>11</sub>	A <sub>11</sub>	A <sub>11</sub>	A <sub>11</sub>	B <sub>10</sub>	B <sub>11</sub>	B <sub>12</sub>	B <sub>13</sub>	
C <sub>20</sub>	C <sub>21</sub>	C <sub>22</sub>	C <sub>23</sub>	A <sub>22</sub>	A <sub>22</sub>	A <sub>22</sub>	A <sub>22</sub>	B <sub>20</sub>	B <sub>21</sub>	B <sub>22</sub>	B <sub>23</sub>	
C <sub>30</sub>	C <sub>31</sub>	C <sub>32</sub>	C <sub>33</sub>	A <sub>33</sub>	A <sub>33</sub>	A <sub>33</sub>	A <sub>33</sub>	B <sub>30</sub>	B <sub>31</sub>	B <sub>32</sub>	B <sub>33</sub>	

## Execution Steps...

C <sub>00</sub>	C <sub>01</sub>	C <sub>02</sub>	C <sub>03</sub>	A <sub>00</sub>	A <sub>01</sub>	A <sub>02</sub>	A <sub>03</sub>	B <sub>10</sub>	B <sub>11</sub>	B <sub>12</sub>	B <sub>13</sub>	Shift of B
C <sub>10</sub>	C <sub>11</sub>	C <sub>12</sub>	C <sub>13</sub>	A <sub>10</sub>	A <sub>11</sub>	A <sub>12</sub>	A <sub>13</sub>	B <sub>20</sub>	B <sub>21</sub>	B <sub>22</sub>	B <sub>23</sub>	
C <sub>20</sub>	C <sub>21</sub>	C <sub>22</sub>	C <sub>23</sub>	A <sub>20</sub>	A <sub>21</sub>	A <sub>22</sub>	A <sub>23</sub>	B <sub>30</sub>	B <sub>31</sub>	B <sub>32</sub>	B <sub>33</sub>	
C <sub>30</sub>	C <sub>31</sub>	C <sub>32</sub>	C <sub>33</sub>	A <sub>30</sub>	A <sub>31</sub>	A <sub>32</sub>	A <sub>33</sub>	B <sub>00</sub>	B <sub>01</sub>	B <sub>02</sub>	B <sub>03</sub>	
C <sub>00</sub>	C <sub>01</sub>	C <sub>02</sub>	C <sub>03</sub>	A <sub>01</sub>	A <sub>01</sub>	A <sub>01</sub>	A <sub>01</sub>	B <sub>10</sub>	B <sub>11</sub>	B <sub>12</sub>	B <sub>13</sub>	Broadcast of A's 2nd diag. (stored in a Separate buffer)
C <sub>10</sub>	C <sub>11</sub>	C <sub>12</sub>	C <sub>13</sub>	A <sub>12</sub>	A <sub>12</sub>	A <sub>12</sub>	A <sub>12</sub>	B <sub>20</sub>	B <sub>21</sub>	B <sub>22</sub>	B <sub>23</sub>	
C <sub>20</sub>	C <sub>21</sub>	C <sub>22</sub>	C <sub>23</sub>	A <sub>23</sub>	A <sub>23</sub>	A <sub>23</sub>	A <sub>23</sub>	B <sub>30</sub>	B <sub>31</sub>	B <sub>32</sub>	B <sub>33</sub>	
C <sub>30</sub>	C <sub>31</sub>	C <sub>32</sub>	C <sub>33</sub>	A <sub>30</sub>	A <sub>30</sub>	A <sub>30</sub>	A <sub>30</sub>	B <sub>00</sub>	B <sub>01</sub>	B <sub>02</sub>	B <sub>03</sub>	
C <sub>00</sub>	C <sub>01</sub>	C <sub>02</sub>	C <sub>03</sub>	A <sub>01</sub>	A <sub>01</sub>	A <sub>01</sub>	A <sub>01</sub>	B <sub>10</sub>	B <sub>11</sub>	B <sub>12</sub>	B <sub>13</sub>	Local computation
C <sub>10</sub>	C <sub>11</sub>	C <sub>12</sub>	C <sub>13</sub>	A <sub>12</sub>	A <sub>12</sub>	A <sub>12</sub>	A <sub>12</sub>	B <sub>20</sub>	B <sub>21</sub>	B <sub>22</sub>	B <sub>23</sub>	
C <sub>20</sub>	C <sub>21</sub>	C <sub>22</sub>	C <sub>23</sub>	A <sub>23</sub>	A <sub>23</sub>	A <sub>23</sub>	A <sub>23</sub>	B <sub>30</sub>	B <sub>31</sub>	B <sub>32</sub>	B <sub>33</sub>	
C <sub>30</sub>	C <sub>31</sub>	C <sub>32</sub>	C <sub>33</sub>	A <sub>30</sub>	A <sub>30</sub>	A <sub>30</sub>	A <sub>30</sub>	B <sub>00</sub>	B <sub>01</sub>	B <sub>02</sub>	B <sub>03</sub>	

## Fox's Algorithm

```
// No initial data movement
for k = 1 to q in parallel
    Broadcast A's kth diagonal
    Local C = C + A*B
    Vertical shift of B
// No final data movement
```

- Again note that there is an additional array to store incoming diagonal block
- This is the array we use in the A\*B multiplication

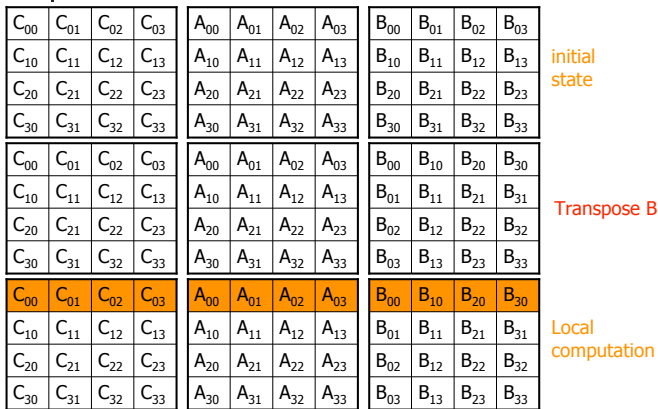
## Performance Analysis

- You'll have to do it in a homework assignment
  - Write pseudo-code of the algorithm in more details
  - Write the performance analysis

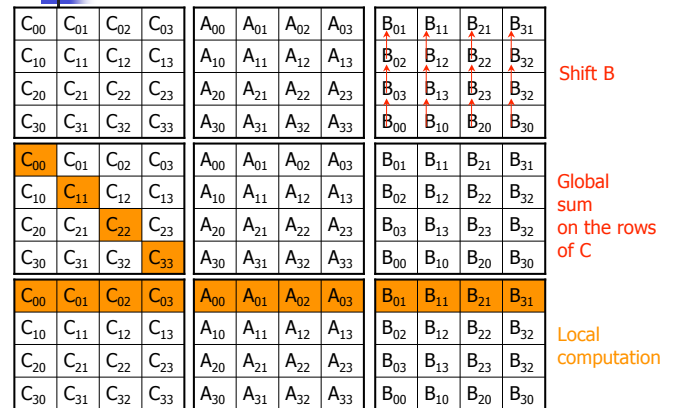
## Snyder's Algorithm (1992)

- More complex than Cannon's or Fox's
- First transposes matrix B
- Uses reduction operations (sums) on the rows of matrix C
- Shifts matrix B

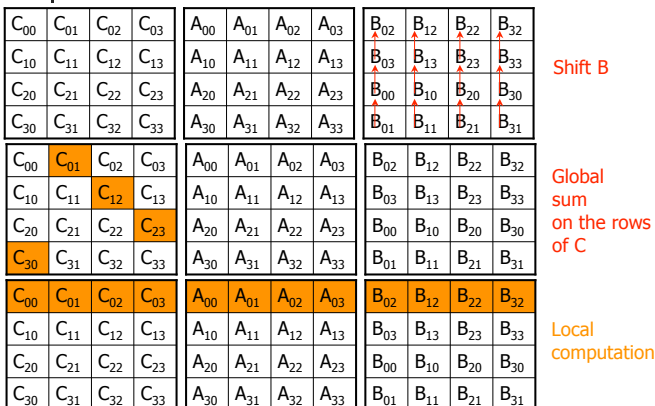
## Execution Steps...



## Execution Steps...



## Execution Steps...



## The Algorithm

```

var A,B,C: array[0..m-1][0..m-1] of real
var bufferC: array[0..m-1][0..m-1] of real
Transpose B
MatrixMultiplyAdd(bufferC, A, B, m)
Vertical shifts of B
For k = 1 to q-1
    Global sum of bufferC on proc rows into  $C_{i,(i+k-1)\%q}$ 
    MatrixMultiplyAdd(bufferC, A, B, m)
    Vertical shift of B
    Global sum of bufferC on proc rows into  $C_{i,(i+k-1)\%q}$ 
    Transpose B
    
```

## Performance Analysis

- The performance analysis isn't fundamentally different than what we've done so far
- But it's a bit cumbersome
- In a homework you'll do the algorithm for transposing a matrix and come up with the performance analysis of this algorithm
  - Just corresponds to the pre- and post-skewing for the Snyder Algorithm


## Which Data Distribution?

- So far we've seen:
  - Block Distributions
  - 1-D Distributions
  - 2-D Distributions
  - Cyclic Distributions
- One may wonder what a good choice is for a data distribution?
- Many people argue that a good "Swiss Army knife" is the "2-D block cyclic distribution"


## The 2-D block cyclic distribution

- Goal: try to have all the advantages of both the horizontal and the vertical 1-D block cyclic distribution
  - Works whichever way the computation “progresses”
    - left-to-right, top-to-bottom, wavefront, etc.
- Consider a number of processors  $p = r * c$ 
  - arranged in a  $r \times c$  matrix
- Consider a 2-D matrix of size  $N \times N$
- Consider a block size  $b$  (which divides  $N$ )

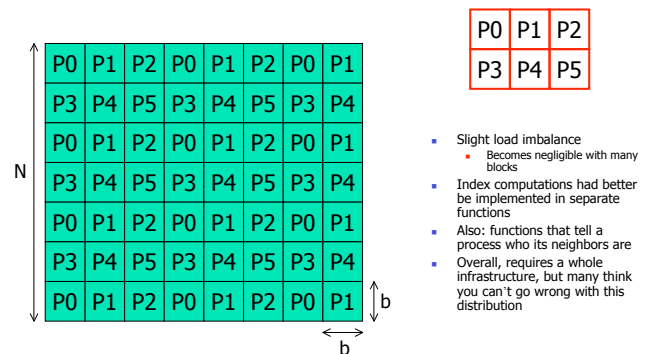
## The 2-D block cyclic distribution



## The 2-D block cyclic distribution



## The 2-D block cyclic distribution



- Slight load imbalance
  - Becomes negligible with many blocks
- Index computations had better be implemented in separate functions
- Also: functions that tell a process who its neighbors are
- Overall, requires a whole infrastructure, but many think you can't go wrong with this distribution

## Conclusion

- All the algorithms we have seen in the semester can be implemented on a 2-D block cyclic distribution
- The code ends up much more complicated
- But one may expect several benefits “for free”
- The ScaLAPACK library recommends to use the 2-D block cyclic distribution
  - Although its routines support all other distributions